

Réunion Louis + Nahuel + Konstantin



Seminar of the NEO team (Inria Sophia Antipolis)

Louis HAUSEUX (Inria Sophia Antipolis, superv. by K. AVRACHENKOV & J. ZERUBIA)

Joint work with Nahuel SOPRANO-LOTO (Inria Paris)

and Konstantin AVRACHENKOV (Inria Sophia Antipolis)

Inria Sophia Antipolis, Lagrange Gris room (L101) – 2.30 p.m. Thursday 16 July 2026



Outline

1. Motivation & the Bayesian framework
2. The Sankararaman & Baccelli model
3. MCMC Dynamics & Performance Bounds
4. Open Questions & Conclusion

01

Motivation & the Bayesian framework



Haplotype Assembly

1. Biological Setup: Two Homologous Chromosomes (Haplotypes)

Humans are diploid: they have two copies of each chromosome (homologous chromosomes).

Haplotype 1
(Truth $\Sigma = +$)


Haplotype 2
(Truth $\Sigma = -$)



Haplotype Assembly

2. High-Throughput Sequencing: DNA Reads

Sequencing generates short, overlapping fragments called **reads** from both chromosomes.






Haplotype Assembly

3. Modeling as a Signed Graph

Nodes = reads Edges = overlaps

Goal: Reconstruct the true haplotype assignments Σ from this observed signed graph.





Motivation: Signed Community Detection

Community Detection on Signed Graphs ($K = 2$)

- ▶ Partition nodes into $K = 2$ communities: $\Sigma_i \in \{+1, -1\}$.
- ▶ Links represent **attractions** (solid) and **repulsions** (dashed).
- ▶ Beyond the classical SBM: a rich framework coupling spatial geometry and community structure!

The Geometric Challenge

- ▶ Spatial embedding (e.g., positions on $\mathbb{Z}$ or $\mathbb{R}^d$).
- ▶ Edge probabilities depend on distance $r = \|x_i - x_j\|$:

$$\mathbb{P}(\text{attractive edge} \mid \text{same community}) = f_{\text{in}}(r)$$

$$\mathbb{P}(\text{attractive edge} \mid \text{different communities}) = f_{\text{out}}(r)$$

- ▶ **Main interest of the model... but also core challenge:** capturing this geometry.




The General Bayesian Model

Definition

1. **Nodes, Edges, Communities:** Nodes $V_n = \{1, \dots, n\}$, edges E_n , communities $\mathcal{L}_K = \{0, \dots, K-1\}$, and set of configurations $\sigma \in \mathcal{C}_n = (\mathcal{L}_K)^{V_n}$.
2. **Features:** $X_n \in \mathcal{X}_n$ according to distribution ν_n .
3. **Prior on configurations:** $\mu_{0,n}(\mathrm{d}\sigma \mid x)$ over $\mathcal{C}_n$.
4. **(Noisy) channel:** $Q_n(\mathrm{d}w \mid x, \sigma) = \prod_{e \in E_n} Q_{e,n}(\mathrm{d}w_e \mid x_i, x_j, \sigma_i, \sigma_j)$.

The joint law is $\mathbb{P}_n(\mathrm{d}x, \mathrm{d}\sigma, \mathrm{d}w) = \nu_n(\mathrm{d}x) \mu_{0,n}(\mathrm{d}\sigma \mid x) Q_n(\mathrm{d}w \mid x, \sigma)$.







Gibbs Model on Signed Graphs

Attractive and Repulsive Edges

Partition potential edges $E = F \dot{\cup} \tilde{F}$:

- ▶ F : **Attractive** (ferromagnetic), weight $W_e > 0$
- ▶ $\tilde{F}$: **Repulsive** (antiferromagnetic), weight $W_e < 0$



Signed Gibbs Measure [1]

Configuration probability $\mu(\sigma) = \frac{1}{Z} e^{-U(\sigma)}$, where energy $U(\sigma)$ penalizes unsatisfied edges:

$$U(\sigma) = \sum_{e=\{i,j\} \in F} W_e \mathbb{I}_{\{\sigma_i \neq \sigma_j\}} + \sum_{e=\{i,j\} \in \tilde{F}} |W_e| \mathbb{I}_{\{\sigma_i = \sigma_j\}}$$

Gibbs Assumption

Special case of the Bayesian framework where the posterior is assumed to be:

$$\mu_{X,W}(\sigma) = \frac{1}{Z} e^{-U(\sigma)}$$



Advantages of the Bayesian Framework

Rigorous Definitions

- ▶ **Algorithm** τ_n : MARKOV kernel mapping observations (X, W) to configurations.
- ▶ **Random Guess**: An algorithm agnostic to observations (relying only on the prior μ_0).
- ▶ **Weak / Partial Recovery**: Beating the best random guess asymptotically.

Posterior Resampling

Replace hidden truth by replica:

$$\Sigma \iff \Sigma^{(1)} \sim \mu(\cdot \mid X, W)$$

Measure of Performance: Overlap

The overlap quotients the HAMMING similarity by label permutations $\mathfrak{S}_K$:

$$\text{ov}_n(\sigma, \sigma') = \max_{\pi \in \mathfrak{S}_K} \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\sigma'_i = \pi(\sigma_i)\}}$$

- ▶ Uniform prior: best random guess score is $1/K$.
- ▶ Weak recovery: $\text{overlap} \geq 1/K + \eta$.



Posterior Resampling (Nishimori Identity)

Core Idea: Replacing Truth by a Replica

Comparing τ to truth $\Sigma \iff$ Comparing τ to replica $\Sigma^{(1)} \sim \mu$

Theorem (Nishimori Identity [2])

For any algorithm τ , posterior replica $\Sigma^{(1)} \sim \mu$, and bounded test function Ψ :

$$\mathbb{E} [\Psi(X, W, \Sigma, \tau)] = \mathbb{E} [\Psi(X, W, \Sigma^{(1)}, \tau)]$$

Corollary (Overlap Identity)

$$\mathbb{P} [\text{ov}(\Sigma, \tau) \geq s] = \mathbb{P} [\text{ov}(\Sigma^{(1)}, \tau) \geq s] \quad (\forall s \in [0, 1])$$

02

The Sankararaman & Baccelli model





The Sankararaman–Bacelli Model (GSBM)

Model Definition [3]

- **Nodes:** Realization X of a homogeneous POISSON point process in an n -volume of $\mathbb{R}^d$ [4].
- **Latent Communities:** Uniform assignments $\Sigma_i \in \{+1, -1\}$ ($K = 2$).
- **Edges:** Connection $e = \{i, j\} \in H$ occurs with probability:

$$\mathbb{P}(e \in H \mid X, \Sigma) = \begin{cases} f_{\text{in}}(r_e) & \text{if } \Sigma_i = \Sigma_j \text{ (attraction),} \\ f_{\text{out}}(r_e) & \text{if } \Sigma_i \neq \Sigma_j \text{ (repulsion),} \end{cases}$$

where $r_e = \|X_i - X_j\|$ and $f_{\text{in}}(r) \geq f_{\text{out}}(r)$.

Goal: Weak Recovery

- Find an algorithm τ_n (independent of Σ_n) achieving:

$$\lim_{n \rightarrow \infty} \mathbb{E}[\text{ov}_n(\Sigma_n, \tau_n)] \geq \frac{1}{2} + \eta \quad (\eta > 0), \quad \text{where } \text{ov}_n(\Sigma_n, \tau_n) = \max_{\pi \in \mathfrak{S}_2} \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\tau_n, i = \pi(\Sigma_n, i)\}}.$$



GSBM in the Bayesian Framework

Sankararaman & Baccelli (GSBM)	Bayesian Framework
Positions X : Realization of a homogeneous POISSON point process Prior : Uniform prior μ_0 on $\Sigma \in \{+1, -1\}^n$ Observables : (H, X) where H is the observed graph	Features $X \sim \nu$: same POISSON law Prior μ_0 : same Observables : (W, X) where W is the channel output
Deterministic construction of W from H: $W_e = \begin{cases} \log \left(\frac{f_{\text{in}}(r_e)}{f_{\text{out}}(r_e)} \right) \geq 0 & \text{if } e \in H \text{ (attractive),} \\ -\log \left(\frac{1 - f_{\text{out}}(r_e)}{1 - f_{\text{in}}(r_e)} \right) \leq 0 & \text{if } e \notin H \text{ (repulsive).} \end{cases}$	
Posterior : $\mathbb{P}(\Sigma = \sigma \mid H, X)$	Posterior : $= \mathbb{P}(\Sigma = \sigma \mid W, X) = \mu(\sigma)$  GIBBS assumption



The Percolation Obstruction

Necessary Condition for Weak Recovery [3]

- Requires a global information flow graph to **percolate** [5] (possess a giant component).

Theorem (Sankararaman & Baccelli [3])

Independent bond percolation with edge probability

$$p_e = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

*yielding a giant component is a **necessary condition** for weak recovery.*

- **Our goal:** Retrieve and extend this result using unified **Bayesian MCMC coupling**.

03

MCMC Dynamics & Performance Bounds






Invariant Coupling

Markov Chain Dynamics as a Source of Coupling


- ▶ We construct the posterior replica $\Sigma^{(1)}$ starting from the hidden truth Σ .
- ▶ Running **exactly one step** of a posterior-invariant MARKOV kernel M initialized at Σ yields a valid replica $\Sigma^{(1)} \sim M(\Sigma, \cdot)$.



Single-Spin Dynamics & Geometric Bottleneck


Single-Spin Update Limitations

- **Geometry makes updates slow:** Single-spin dynamics (GLAUBER / METROPOLIS-HASTINGS [6]) update one node at a time.
- **Contradictory Clusters:** Two internally valid (highly satisfied) but mutually contradictory clusters are locked.
- **High Energy Barrier:** Flipping a single boundary node violates internal edges, creating a local energy barrier.
- **Conclusion:** Need for geometry $\implies$ **Cluster Dynamics** (updating whole regions).





One Step of Signed Swendsen–Wang [7, 8]



- Edges in bold represent satisfied interactions that are randomly frozen.
- Clusters are then flipped independently.



Swendsen–Wang Signed Cluster Dynamics

Swendsen–Wang (SW) Signed Cluster Dynamics [7, 8]

- ▶ Update whole clusters of vertices simultaneously to speed up mixing.
- ▶ **Freezing Step:** For each edge $e = \{i, j\}$:
 - If e is satisfied ($W_e > 0$ and $\sigma_i = \sigma_j$, or $W_e < 0$ and $\sigma_i \neq \sigma_j$), freeze it with probability:
$$p_e^{\text{frozen}} = 1 - e^{-|W_e|}.$$
 - If e is unsatisfied, do not freeze it ($p_e^{\text{frozen}} = 0$).
- ▶ **Recoloring Step:** Flip each frozen cluster independently with probability $1/2$.



High-Order Triangular Dynamics

The Frustration Obstruction

- ▶ **Frustration:** $\nexists \sigma$ satisfying all edges (conflicting constraints).
- ▶ **Excessive Freezing:** In frustrated configurations, standard SW satisfies and freezes too many edges independently, blocking/freezing the MCMC dynamics.
- ▶ **Higher-Order Correlation:** Group interactions into **triangles** to correlate freezing and **freeze** as little as possible.

Triangular Cluster Dynamics [9]

- ▶ **Energy Decomposition:** Distribute edge energy U over triangles: $U(\sigma) = \sum_{t \in \mathcal{T}} U_t(\sigma)$.
- ▶ **Freezing Mechanism:** Apply a joint freezing mechanism on the triangles (rather than independent edges) to keep cluster sizes active.



Triangular Dynamics Freezing Rules

$\triangle$ Attractive Triangle (Isotropic)					
Configuration	$\emptyset$	$\{ab\}$	$\{ac\}$	$\{bc\}$	$\{ab, ac, bc\}$
$(+, +, +)$	b_u	0	0	0	a_u
Other	1	0	0	0	0

$\triangleleft \triangle$ Repulsive Triangle (Isotropic)					
Configuration	$\emptyset$	$\{ab\}$	$\{ac\}$	$\{bc\}$	$\{ab, ac, bc\}$
$(+, +, +)$	1	0	0	0	0
$(+, +, -)$	b_u	0	$a_u/2$	$a_u/2$	0
$(+, -, +)$	b_u	$a_u/2$	0	$a_u/2$	0
$(+, -, -)$	b_u	$a_u/2$	$a_u/2$	0	0

Table: Freezing rules for triangular dynamics ($a_u = 1 - e^{-2u}$, $b_u = e^{-2u}$), where u is the energy transferred by each edge.






General Swendsen–Wang Cluster Dynamics: Framework

General Framework [10]

Suppose the Hamiltonian decomposes as: $U(\sigma) = U_0(\sigma) + \sum_{b \in \mathcal{B}} U_b(\sigma)$, where each $b \in \mathcal{B}$ is a **bond** (e.g., edge, triangle, clique).

1. **Freezing Step:** For each bond $b \in \mathcal{B}$, freeze a sub-bond $\omega_b \subseteq b$ with probability $P_b^\sigma(\omega_b)$.
2. **Recoloring Step:** Recolor each frozen cluster independently according to a distribution preserving the prior $\mu_0(\sigma) \propto e^{-U_0(\sigma)}$.





General Swendsen–Wang Cluster Dynamics: Invariance

Definition (Local Reversibility)

For two configurations σ, σ' and a bond b , we define the compatibility probability:

$$P_b^{\sigma \rightarrow \sigma'}(\omega_b) = \begin{cases} P_b^\sigma(\omega_b), & \text{if } P_b^{\sigma'}(\omega_b) > 0, \\ 0, & \text{otherwise.} \end{cases}$$

The freezing rule is **locally reversible** if for all $b, \omega_b, \sigma, \sigma'$:

$$P_b^{\sigma \rightarrow \sigma'}(\omega_b) = e^{-(U_b(\sigma') - U_b(\sigma))} P_b^{\sigma' \rightarrow \sigma}(\omega_b)$$

Theorem (Gibbs Invariance [10])

If freezing choices are independent, locally reversible, and $P_b^\sigma(\emptyset) > 0$ for all b, σ :

1. The global cluster dynamics is reversible with respect to $\mu(\sigma) \propto e^{-U(\sigma)}$.
2. The condition $P_b^\sigma(\emptyset) > 0$ ensures the **irreducibility** of the MARKOV chain.

Percolation as a Necessary Condition for Recovery

Theorem (Percolation Bound on Overlap [10])

Under the uniform prior μ_0 , if frozen clusters are recolored independently and uniformly, then for the one-step posterior replica $\Sigma_n^{(1)}$, any algorithm τ_n , and any $\eta > 0$:

$$\lim_{n \rightarrow \infty} \mathbb{P} \left[\text{ov}_n(\Sigma_n^{(1)}, \tau_n) \geq \frac{1}{K} + \frac{K-1}{K}(\theta^{\max} + \eta) \right] = 0$$

where $\theta^{\max}$ is the asymptotic fraction of nodes in giant frozen components (which bounds the maximum recoverable fraction of communities).

Corollary (Percolation is Necessary)

*If there is no percolation ($\theta^{\max} = 0$), the maximum overlap is bounded by $1/K$ (random guess): weak recovery is impossible. Hence, **percolation is a necessary condition for weak recovery**.*

Key Intuition

Small clusters are recolored independently, washing out their contribution. Only giant (macroscopic) components can carry global community information.

Coupling & the Sankararaman–Baccelli Bound

One Step of Classical Edge-by-Edge Swendsen–Wang

Let's run one step of standard SWENDSEN–WANG (edges) initialized at the hidden truth Σ . For any edge e , the probability that it is frozen is:

► **Case 1** ($\Sigma_i = \Sigma_j$):

$$\mathbb{P}(\text{frozen}) = f_{\text{in}}(r_e) \left(1 - e^{-|W_e|}\right) = f_{\text{in}}(r_e) \left(1 - \frac{f_{\text{out}}(r_e)}{f_{\text{in}}(r_e)}\right) = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

► **Case 2** ($\Sigma_i \neq \Sigma_j$):

$$\mathbb{P}(\text{frozen}) = (1 - f_{\text{out}}(r_e)) \left(1 - e^{-|W_e|}\right) = (1 - f_{\text{out}}(r_e)) \left(1 - \frac{1 - f_{\text{in}}(r_e)}{1 - f_{\text{out}}(r_e)}\right) = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

Corollary (Retrieving the S&B Bound [3])

*Under the classical edge SW dynamics, the frozen graph is **exactly** an independent bond percolation with edge probability:*

$$p_e = f_{\text{in}}(r_e) - f_{\text{out}}(r_e).$$



Strict Improvement of the S&B bound (on the triangular grid)

Homogeneous Model ($f_{\text{in}} = p$, $f_{\text{out}} = 1 - p$)

On the triangular grid, we select a family of independent white triangles (like a chessboard).

- **Edge Dynamics (Standard SW):** Frozen graph is independent edge percolation with parameter $p_{\text{edge}} = 2p - 1$. Critical threshold:

$$p_c^{\text{edge}} = \frac{1}{2} + \sin\left(\frac{\pi}{18}\right) \approx 0.673 \dots$$

- **Triangular Dynamics:** Frozen graph yields correlated triangle percolation. Critical threshold:


$$p_c^{\triangle} = \frac{7}{4} - \frac{\sqrt{17}}{4} \approx 0.7192 \dots$$

Corollary (Strictly Better Impossibility Bound [10])

Since $p_c^{\text{edge}} < p_c^{\triangle}$, for any $p \in [0.673 \dots, 0.7192 \dots)$, the standard S&B bound cannot conclude (edge dynamics percolate), whereas triangular dynamics do not, proving that weak recovery is impossible.



Empirical Results: Subcritical vs. Supercritical ($p = 0.70$)



Phase Space


$(p_c^{\text{edge}} \approx 0.673 < 0.70 < p_c^{\Delta} \approx 0.7192)$

- ▶ **Edges (standard SW):** supercritical ($\theta^{\max} > 0$), largest component size (LCC) reaches ≈ 0.9 (right plot).
- ▶ **Triangles / Split weight:** subcritical ($\theta^{\max} = 0$), LCC remains near 0.

Observed Behavior

- ▶ Only standard edge SW achieves giant components.
- ▶ Overlap remains at the baseline noise level (around 0.002) for triangles and split weights due to lack of percolation (left plot).

Empirical Results: Onset of Percolation ($p = 0.72$)



Transition Regime

($p = 0.72 > p_c^\Delta \approx 0.7192$)


- ▶ All three dynamics are now **supercritical**.
- ▶ LCC grows for all: edges reach ≈ 0.9 , split weights reach ≈ 0.85 , and white triangles reach ≈ 0.55 (right plot).

Observed Behavior

- ▶ Giant components have formed, but the overlap remains low (below 0.003) for all dynamics (left plot).
- ▶ **No weak recovery yet:** the true weak recovery threshold is likely strictly above the percolation threshold (percolation is only a necessary condition).



Empirical Results: Weak Recovery Achieved ($p = 0.80$)



Supercritical Regime ($p = 0.80 \gg p_c^\Delta$)

- ▶ Giant components are well-established (right plot), with LCC close to 0.95 for edges and split weights, and 0.90 for white triangles.
- ▶ Fluctuations are small and components are stable.

Observed Behavior

- ▶ Weak recovery is successfully achieved for all algorithms.
- ▶ The overlap climbs significantly above the noise level (left plot), peaking at ≈ 0.010 for white triangles and ≈ 0.006 for split weights.



Target Gibbs Law and Exponential Clocks

Let $E = F \dot{\cup} \tilde{F}$. An edge $e = \{i, j\}$ is **satisfied** when

$$\sigma_i = \sigma_j \quad (e \in F), \quad \sigma_i \neq \sigma_j \quad (e \in \tilde{F}).$$

Thus, up to an additive constant,

$$U(\sigma) = \sum_{e \in E} |W_e| \mathbb{I}_{\{e \text{ is not satisfied}\}},$$

$$\mu(\sigma) \propto \mu_0(\sigma) e^{-U(\sigma)}.$$


For each satisfied edge, independently,

$$\xi_e \sim \text{Exp}(|W_e|),$$

while $\xi_e = +\infty$ for an unsatisfied edge. At level t , retain precisely the links with $\xi_e \leq t$.

Classical Swendsen–Wang is the cut $t = 1$

$$\mathbb{P}(\xi_e \leq 1 \mid e \text{ satisfied}) = 1 - e^{-|W_e|}.$$



Continuous clocks imply that merge times are distinct almost surely.



The Auxiliary Variable Is the Component Dendrogram up to 1

Define the graph and its component partition at every level by

$$G_t = (V, \{e \in E : \xi_e \leq t\}), \quad \Pi_t = \text{cc}(G_t), \quad \boxed{D_{\leq 1} = (\Pi_t)_{0 \leq t \leq 1}.}$$

Finite description of $D_{\leq 1}$


- ▶ the final components (the roots) Π_1 ;
- ▶ for each merger $u : C = C_1 \dot{\cup} C_2$, the two children C_1, C_2 and its height $\beta_u \leq 1$.

Children are unordered. The support is

$$0 < \beta_v < \beta_u \leq 1 \quad \text{whenever } v \text{ is below } u.$$

The reference measure is counting measure on these forest topologies times Lebesgue measure on their admissible heights.


----- $t = 1$



Conditioned data at a node. At u , they are exactly (C_1, C_2, β_u) ; the likelihood below sums over every graph link between C_1 and C_2 .



A Minimum Spanning Forest Computes All Those Partitions



Exact pruning property

For every $0 \leq t \leq 1$,

$$\begin{aligned}\Pi_t &= \text{cc}\{e : \xi_e \leq t\} \\ &= \text{cc}\{e \in \text{MSF} : \xi_e \leq t\}.\end{aligned}$$

The MSF is therefore a computational representation of the same nested partitions. The auxiliary variable used below remains $D_{\leq 1}$.


Kruskal gives the hierarchy; the probability law is derived from the full set of graph links.

At One Merge, All Links Across the Cut Race

For $u : C = C_1 \dot{\cup} C_2$, use the whole cut

$$E_u = \{e = \{i, j\} \in E : i \in C_1, j \in C_2\},$$

$$\Lambda_u(\sigma) = \sum_{e \in E_u} |W_e| \mathbb{I}_{\{e \text{ satisfied by } \sigma\}}.$$



$$\beta_u = \min_{\substack{e \in E_u \\ e \text{ satisfied}}} \xi_e.$$

The possible first bell is summed over all satisfied $e \in E_u$:

$$\sum_{e \in E_u} |W_e| \mathbb{I}_{\{e \text{ satisfied}\}} e^{-\beta_u \Lambda_u(\sigma)} = \Lambda_u(\sigma) e^{-\beta_u \Lambda_u(\sigma)}.$$



Exact Density of the Censored Dendrogram

Fix a candidate forest topology. Assign each graph edge to

- ▶ the unique split E_u given by the first common cluster of its endpoints; or
- ▶ E_∞ , if its endpoints lie in distinct roots of Π_1 .

Hence

$$E = \left(\dot{\bigcup}_{u \in \text{Int}(D_{\leq 1})} E_u \right) \dot{\cup} E_\infty, \quad \Lambda_\infty(\sigma) = \sum_{e \in E_\infty} |W_e| \mathbb{I}_{\{e \text{ satisfied}\}}.$$

Internal node u

The first clock in E_u has density

$$\Lambda_u(\sigma) e^{-\beta_u \Lambda_u(\sigma)} d\beta_u.$$

Between final roots

No satisfied clock in E_∞ rings before 1:

$$\mathbb{P}\left(\min_{e \in E_\infty} \xi_e > 1 \mid \sigma\right) = e^{-\Lambda_\infty(\sigma)}.$$

The clock families in this disjoint partition are independent. Therefore

$$\mathbb{P}(dD_{\leq 1} \mid \sigma) = \mathbb{I}_{\{D_{\leq 1} \text{ admissible}\}} e^{-\Lambda_\infty(\sigma)} \prod_{u \in \text{Int}(D_{\leq 1})} \Lambda_u(\sigma) e^{-\beta_u \Lambda_u(\sigma)} d\beta_u.$$

A topology with $\Lambda_u(\sigma) = 0$ has zero density, as it should.

The Edwards–Sokal Lift and Its Conditional Law

Define the joint law on spin–dendrogram pairs by

$$\nu(d\sigma, dD) = \frac{1}{Z} \mu_0(d\sigma) e^{-U(\sigma)} \mathbb{P}(dD_{\leq 1} \mid \sigma).$$

Since $\mathbb{P}(\cdot \mid \sigma)$ is a probability law, the σ -marginal of ν is exactly the target Gibbs law μ .

Every edge is in one E_u or in E_∞ , so

$$\sum_{e \in E} |W_e| \mathbb{I}_{\{e \text{ satisfied}\}} = \sum_u \Lambda_u(\sigma) + \Lambda_\infty(\sigma).$$

Moreover

$$e^{-U(\sigma)} = e^{-\sum_{e \in E} |W_e| \sum_u \Lambda_u(\sigma) + \Lambda_\infty(\sigma)}.$$

The Λ_∞ terms cancel exactly at the cut 1.

Thus, on the support of ν ,

Conditional law used by every flip

$$\nu(\sigma \mid D) \propto \mu_0(\sigma) \prod_u \Lambda_u(\sigma) e^{(1-\beta_u)\Lambda_u(\sigma)}.$$



At a Split There Are Four Joint Events

For $u : C = C_1 \dot{\cup} C_2$, let σ^{ab} flip C_1 when $a = 1$ and C_2 when $b = 1$. In the binary signed model,

(a, b)	C_1	C_2	$\Lambda_u(\sigma^{ab})$
$(0, 0)$	keep	keep	$\Lambda_u(\sigma)$
$(1, 0)$	flip	keep	$\sum_{e \in E_u} W_e - \Lambda_u(\sigma)$
$(0, 1)$	keep	flip	$\sum_{e \in E_u} W_e - \Lambda_u(\sigma)$
$(1, 1)$	flip	flip	$\Lambda_u(\sigma)$

One four-state draw, not two independent coins

The current split only distinguishes equal versus opposite flip parity. The prior μ_0 and the ancestor cuts generally distinguish all four events.

If u is a root and μ_0 is uniform, set $A_u = \sum_{e \in E_u} |W_e|$. Then the fully explicit local probabilities are

$$p_{00} = p_{11} = \frac{\Lambda_u e^{(1-\beta_u)\Lambda_u}}{2[\Lambda_u e^{(1-\beta_u)\Lambda_u} + (A_u - \Lambda_u)e^{(1-\beta_u)(A_u - \Lambda_u)}]},$$

$$p_{10} = p_{01} = \frac{(A_u - \Lambda_u)e^{(1-\beta_u)(A_u - \Lambda_u)}}{2[\Lambda_u e^{(1-\beta_u)\Lambda_u} + (A_u - \Lambda_u)e^{(1-\beta_u)(A_u - \Lambda_u)}]}.$$

The Exact Four Probabilities at an Arbitrary Level

For the general node u , factors below u and on disjoint branches are unchanged. The exact heat bath is therefore

$$p_u(a, b \mid \sigma, D) = \frac{\mu_0(\sigma^{ab}) \prod_{v \succeq u} \left[\Lambda_v(\sigma^{ab}) e^{(1-\beta_v)\Lambda_v(\sigma^{ab})} \right]}{\sum_{c,d \in \{0,1\}} \mu_0(\sigma^{cd}) \prod_{v \succeq u} \left[\Lambda_v(\sigma^{cd}) e^{(1-\beta_v)\Lambda_v(\sigma^{cd})} \right]}.$$

Here $v \succeq u$ means $v = u$ or v is an ancestor of u , and

$$\Lambda_v(\sigma^{ab}) = \sum_{e \in E_v} |W_e| \mathbb{I}_{\{e \text{ satisfied by } \sigma^{ab}\}}.$$

Where the clock minimum enters

The sampled height β_v is the minimum crossing clock used to construct D . It is then held fixed and appears explicitly in every factor $e^{(1-\beta_v)\Lambda_v}$.

Why a purely local rule is wrong

Flipping C_1 or C_2 changes links crossing ancestor splits. Omitting μ_0 or those ancestor factors generally breaks detailed balance.

Detailed Balance for Each Four-State Heat Bath

Fix D and u , and let

$$\mathcal{O}_u(\sigma) = \{\sigma^{00}, \sigma^{10}, \sigma^{01}, \sigma^{11}\}.$$

This orbit is the same when viewed from any $x \in \mathcal{O}_u(\sigma)$. The previous rule can be written

$$K_u^D(x, y) = \mathbb{I}_{\{y \in \mathcal{O}_u(x)\}} \frac{\nu(y \mid D)}{\nu(\mathcal{O}_u(x) \mid D)}.$$

Detailed balance, line by line

For x, y in the same orbit,

$$\nu(x \mid D) K_u^D(x, y) = \frac{\nu(x \mid D) \nu(y \mid D)}{\nu(\mathcal{O}_u \mid D)} = \nu(y \mid D) K_u^D(y, x).$$

If x, y are in different orbits, both sides are zero.

Hence K_u^D is reversible and preserves the *same* conditional law $\nu(\cdot \mid D)$. On the support of ν , the orbit denominator is positive; candidates with zero conditional weight simply receive probability zero.



Final Components Need Their Own Two-State Heat Bath

Internal nodes alone miss a singleton root. For every final component $R \in \Pi_1$, also consider the two states σ and σ^R , where all spins in R are flipped.

A root flip preserves every internal Λ_u , while the between-root term has already cancelled.

Therefore

$$\frac{\nu(\sigma^R \mid D)}{\nu(\sigma \mid D)} = \frac{\mu_0(\sigma^R)}{\mu_0(\sigma)},$$

$$\mathbb{P}(\text{flip } R \mid \sigma, D) = \frac{\mu_0(\sigma^R)}{\mu_0(\sigma) + \mu_0(\sigma^R)}.$$

Uniform prior

Each final component, including a singleton, is flipped with probability 1/2: this is precisely the usual SWENDSEN–WANG recolouring at the cut 1.

The internal four-state moves are the additional hierarchical moves.

Why this matters

With two vertices in an unsatisfied state, D has two singleton roots. Without root moves, the proposed chain would be the identity and could not converge from that state.



From Conditional Balance to Gibbs Balance

Let $\mathcal{L}(D)$ be any list of root moves and internal-node moves chosen from the full D before the spins are updated.

Stationary sweep

Every deterministic composition of these heat baths preserves $\nu(\cdot \mid D)$:

$$\nu(\cdot \mid D)K_D = \nu(\cdot \mid D).$$

A deterministic sweep need not itself be reversible.
For such a reversible conditional sweep,

Reversible sweep

Choose the order of $\mathcal{L}(D)$ uniformly at random. Every order and its reverse have the same probability, hence

$$K_D^* = K_D \quad \text{in } L^2(\nu(\cdot \mid D)).$$

$$\begin{aligned} \mu(dx)P(x, dy) &= \int \mu(dx) \mathbb{P}(dD \mid x) K_D(x, dy) \\ &= \int \nu(dD) \nu(dx \mid D) K_D(x, dy), \end{aligned}$$

This is symmetric in x, y : the marginal hierarchical kernel satisfies detailed balance for μ .



Stopping the Updates and Truncating D Are Different

Full $D_{\leq 1}$ is kept

Choose from D —before any flip—a fixed depth, threshold, subtree, or prefix of the traversal. The resulting composition still preserves $\nu(\cdot \mid D)$, hence its marginal preserves μ .

An order or stopping rule that reacts to the spins produced during the sweep needs a separate invariance proof.

Exact conclusion: a full sweep, or any D -measurable partial sweep, has Gibbs stationary; convergence requires a separate ergodicity argument.

Only $D_{\leq t}$ is observed, $t < 1$

This is a different auxiliary variable. If $\Lambda_{\infty}^t(\sigma)$ is the satisfied weight between components of Π_t , then

$$\nu(\sigma \mid D_{\leq t}) \propto \mu_0(\sigma) e^{(1-t)\Lambda_{\infty}^t(\sigma)} \times \prod_{\beta_u \leq t} \Lambda_u(\sigma) e^{(1-\beta_u)\Lambda_u(\sigma)}.$$

The residual boundary factor disappears only at $t = 1$.

04

Open Questions & Conclusion





Open Questions

Limits of the One-Replica Coupling

- ▶ Our framework constructs **one** replica $\Sigma_n^{(1)}$ using MCMC cluster dynamics.
- ▶ This only yields an impossibility bound (**necessary condition** for recovery).

Towards a Necessary and Sufficient Condition

- ▶ Achieving a **necessary and sufficient** condition requires a **two-replica** framework.
- ▶ Open: How to define and analyze two-replica couplings under MCMC?

Towards a Practical Community Detection Algorithm

Limits of Direct Gibbs Sampling

- ▶ Direct posterior sampling does not perform well in practice.
- ▶ Likely due to slow mixing of the underlying MARKOV chain.

Estimating the Correlation Matrix

- ▶ Estimating the correlation matrix $\mathbb{P}(\sigma_i = \sigma_j)$ yields much better results.
- ▶ **Future work:** Understand why correlation estimators are more robust.



Distance Adapted to Cluster Dynamics

Hamming Similarity on the Quotient Space

- ▶ Standard overlap is a **HAMMING** similarity quotiented by label permutations:

$$\text{ov}_n(\sigma, \sigma') = \max_{\pi \in \mathfrak{S}_K} \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\sigma'_i = \pi(\sigma_i)\}}.$$

Markov Transition-Adapted Distance

- ▶ Consider distances reflecting transitions of the cluster dynamics.
- ▶ For standard SW, this corresponds to the **symmetric difference** on satisfied edges:

$$d_{\text{sym}}(\sigma, \sigma') = \frac{1}{2} \sum_{e=\{i,j\}} |W_e| \left| \mathbb{I}_{\{\sigma \text{ satisfies } e\}} - \mathbb{I}_{\{\sigma' \text{ satisfies } e\}} \right|$$

- ▶ Establishes links to **spectral graph analysis** (e.g., LAPLACIAN eigenvalues).

Application: 3-SAT Solving via Signed Graphs

The 3-SAT Problem



Find a truth assignment satisfying a set of clauses, e.g.,

$$\mathcal{C} = x_i \vee \neg x_j \vee x_k$$

Clause Encoding

We introduce a reference node **T** (True):

- ▶ **Triangle** (literals $x_i, \neg x_j, x_k$): edge signs opposite to the literal product.
- ▶ **Tetrahedron (with T)**: edge signs conforming to the literal signs.



Bibliography I

- [1] Jörg Reichardt and Stefan Bornholdt. “Statistical mechanics of community detection”. In: *Phys. Rev. E* 74 (1 July 2006), p. 016110. DOI: [10.1103/PhysRevE.74.016110](https://doi.org/10.1103/PhysRevE.74.016110).
- [2] Hidetoshi Nishimori. “Exact results and critical properties of the Ising model with competing interactions”. In: *Journal of Physics C: Solid State Physics* 13.21 (1980), pp. 4071–4076. DOI: [10.1088/0022-3719/13/21/012](https://doi.org/10.1088/0022-3719/13/21/012).
- [3] Abishek Sankararaman and François Baccelli. “Community Detection on Euclidean Random Graphs”. In: *Proceedings of the Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)* (2018), pp. 2181–2200. DOI: [10.1137/1.9781611975031.142](https://doi.org/10.1137/1.9781611975031.142).
- [4] François Baccelli and Bartomiej Baszczyszyn. *Stochastic Geometry and Wireless Networks, Volume I - Theory*. Vol. 1. Foundations and Trends in Networking. NoW Publishers, 2009, p. 150. DOI: [10.1561/13000000006](https://doi.org/10.1561/13000000006).
- [5] Mathew Penrose. *Random Geometric Graphs*. Vol. 5. Oxford Studies in Probability. Oxford University Press, 2003. DOI: [10.1093/acprof:oso/9780198506263.001.0001](https://doi.org/10.1093/acprof:oso/9780198506263.001.0001).
- [6] Persi Diaconis. “The Markov chain Monte Carlo revolution”. In: *Bulletin of the American Mathematical Society* 46.2 (2009), pp. 179–205. DOI: [10.1090/S0273-0979-08-01238-X](https://doi.org/10.1090/S0273-0979-08-01238-X).

Bibliography II

- [7] Robert H. Swendsen and Jian-Sheng Wang. “Nonuniversal critical dynamics in Monte Carlo simulations”. In: *Phys. Rev. Lett.* 58 (2 Jan. 1987), pp. 86–88. DOI: [10.1103/PhysRevLett.58.86](https://doi.org/10.1103/PhysRevLett.58.86).
- [8] Robert G. Edwards and Alan D. Sokal. “Generalization of the Fortuin-Kasteleyn-Swendsen-Wang representation and Monte Carlo algorithm”. In: *Phys. Rev. D* 38 (6 Sept. 1988), pp. 2009–2012. DOI: [10.1103/PhysRevD.38.2009](https://doi.org/10.1103/PhysRevD.38.2009).
- [9] Louis Hauseux, Nahuel Soprano-Loto, and Konstantin Avrachenkov. “Higher-order Monte Carlo cluster dynamics for community detection in Euclidean graphs”. In: *Asilomar Conference on Signals, Systems, and Computers 2025*. Pacific Grove (CA), United States, Oct. 2025. URL: <https://inria.hal.science/hal-05267074>.
- [10] Louis Hauseux, Nahuel Soprano-Loto, and Konstantin Avrachenkov. *A Bayesian Framework for Community Detection on Signed Graphs. Tighter Bounds with Higher-Order Interaction Markov Chain Monte Carlo Cluster Dynamics*. Article en cours de rédaction. 2026+.

Merci.

