

Community Detection on Signed Graphs: A Bayesian MCMC Coupling Approach



Seminar of the MathNet team (Inria Paris)

Louis HAUSEUX (Inria Sophia Antipolis, superv. by K. AVRACHENKOV & J. ZERUBIA)

Joint work with Nahuel SOPRANO-LOTO (Inria Paris)

and Konstantin AVRACHENKOV (Inria Sophia Antipolis)

Inria Paris, Room Paul ERDÖS (A115) – 2 p.m. Monday 15 June 2026



Outline

1. Motivation & the Sankararaman–Baccelli Model
2. The Bayesian Framework & Coupling
3. MCMC Dynamics & Performance Bounds
4. Open Questions & Conclusion

01

Motivation & the Sankararaman–Baccelli Model



Haplotype Assembly

1. Biological Setup: Two Homologous Chromosomes (Haplotypes)

Humans are diploid: they have two copies of each chromosome (homologous chromosomes).

Haplotype 1
(Truth $\Sigma = +$)


Haplotype 2
(Truth $\Sigma = -$)



Haplotype Assembly

2. High-Throughput Sequencing: DNA Reads

Sequencing generates short, overlapping fragments called **reads** from both chromosomes.






Haplotype Assembly

3. Modeling as a Signed Graph

Nodes = reads Edges = overlaps

Goal: Reconstruct the true haplotype assignments Σ from this observed signed graph.





Motivation: Signed Community Detection

Community Detection on Signed Graphs ($K = 2$)

- ▶ We aim to partition nodes into $K = 2$ communities: $\Sigma_i \in \{+1, -1\}$.
- ▶ Observed links represent **attractions** (solid lines) and **repulsions** (dashed lines).
- ▶ Under the hood, this is a **Stochastic Block Model** (SBM) setup.

The Geometric Challenge

- ▶ Nodes are embedded in space (e.g., genomic positions on $\mathbb{Z}$ or $\mathbb{R}^d$).
- ▶ Edge probabilities depend directly on the distance $r = \|x_i - x_j\|$:

$$\mathbb{P}(\text{attractive edge} \mid \text{same community}) = f_{\text{in}}(r)$$

$$\mathbb{P}(\text{repulsive edge} \mid \text{different communities}) = f_{\text{out}}(r)$$

- ▶ **All the difficulty (and the interest!) comes from this geometric dependency.**



The Sankararaman–Bacelli Model (GSBM)

Model Definition [1]

- ▶ **Spatial Nodes** V_n : Hom. POISSON PP $\mathcal{X}$ in $\mathbb{R}^d$ [2] restrict. to a n -volume window.
- ▶ **Latent Communities**: Uniform assignments $\Sigma_i \in \{+1, -1\}$ ($K = 2$).
- ▶ **Geometric Overlaps**: Observed connection $e = \{i, j\} \in H$ with probability:

$$\mathbb{P}(e \in H \mid X, \Sigma) = \begin{cases} f_{\text{in}}(r_e) & \text{if } \Sigma_i = \Sigma_j \text{ (attraction),} \\ f_{\text{out}}(r_e) & \text{if } \Sigma_i \neq \Sigma_j \text{ (repulsion),} \end{cases}$$

with distance $r_e = \|X_i - X_j\|$ and connecting functions $f_{\text{in}}(r) \geq f_{\text{out}}(r)$.

Goal: Weak Recovery

- ▶ Find an algorithm τ (independent of Σ) reconstructing assignments:

$$\lim_{n \rightarrow \infty} \mathbb{E}[\text{ov}_n(\Sigma, \tau)] \geq \frac{1}{2} + \eta \quad \text{for some } \eta > 0,$$

using the **Community Overlap**: $\text{ov}_n(\Sigma, \tau) = \max_{\pi \in \mathfrak{S}_2} \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\tau_i = \pi(\Sigma_i)\}}.$



Gibbs Model on Signed Graphs

Attractive and Repulsive Edges

We partition the potential edges $E = F \dot{\cup} \tilde{F}$:

- ▶ F : **Attractive** (ferromagnetic) edges, with weight $W_e > 0$
- ▶ $\tilde{F}$: **Repulsive** (antiferromagnetic) edges, with weight $W_e < 0$

Signed Gibbs Measure [3]

The probability of a label configuration $\sigma \in \{+1, -1\}^{V_n}$ is $\mu(\sigma) = \frac{1}{Z} e^{-U(\sigma)}$, where the energy $U(\sigma)$ penalizes unsatisfied edges:

$$U(\sigma) = \sum_{e=\{i,j\} \in F} W_e \mathbb{I}_{\{\sigma_i \neq \sigma_j\}} + \sum_{e=\{i,j\} \in \tilde{F}} |W_e| \mathbb{I}_{\{\sigma_i = \sigma_j\}}$$

The magnitudes $|W_e|$ determine the strength of each constraint.



From GSBM to the Gibbs Posterior

Conditioning on the Point Cloud

- ▶ With fixed spatial positions X , all pairwise distances $r_e = \|X_i - X_j\|$ are known.
- ▶ Under a uniform prior on Σ , the **posterior** distribution given the observed graph H is:

$$\mathbb{P}(\Sigma = \sigma \mid H, X) \propto \mathbb{P}(H \mid X, \sigma)$$

Log-Likelihood as Signed Gibbs Energy

- ▶ Choosing weights $W_e(H)$ as below:

$$W_e(H) = \begin{cases} \log \left(\frac{f_{\text{in}}(r_e)}{f_{\text{out}}(r_e)} \right) > 0 & \text{if } e \in H \text{ (attractive } \in F), \\ -\log \left(\frac{1 - f_{\text{out}}(r_e)}{1 - f_{\text{in}}(r_e)} \right) < 0 & \text{if } e \notin H \text{ (repulsive } \in \tilde{F}). \end{cases}$$

- ▶ ... yields the exact correspondence:

$$\mathbb{P}(H \mid X, \sigma) \propto e^{-U(\sigma)}$$



From GSBM to the Gibbs Posterior

Conditioning on the Point Cloud

- ▶ With fixed spatial positions X , all pairwise distances $r_e = \|X_i - X_j\|$ are known.
- ▶ Under a uniform prior on Σ , the **posterior** distribution given the observed graph H is:

$$\mathbb{P}(\Sigma = \sigma \mid H, X) = \mu(\sigma) \propto \mathbb{P}(H \mid X, \sigma)$$

Log-Likelihood as Signed Gibbs Energy

- ▶ Choosing weights $W_e(H)$ as below:

$$W_e(H) = \begin{cases} \log \left(\frac{f_{\text{in}}(r_e)}{f_{\text{out}}(r_e)} \right) > 0 & \text{if } e \in H \text{ (attractive } \in F), \\ -\log \left(\frac{1 - f_{\text{out}}(r_e)}{1 - f_{\text{in}}(r_e)} \right) < 0 & \text{if } e \notin H \text{ (repulsive } \in \tilde{F}). \end{cases}$$

- ▶ ... yields the exact correspondence:

$$\mathbb{P}(H \mid X, \sigma) \propto e^{-U(\sigma)}$$



The Percolation Obstruction

Necessary Condition for Weak Recovery [1]

- ▶ Fragmented graphs (small, isolated components) prevent global reconstruction.
- ▶ Global information flow requires the graph to **percolate** [4] (possess a giant component).

Theorem (Sankararaman & Baccelli [1])

Independent bond percolation with edge probability

$$p_e = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

*yielding a giant component is a **necessary condition** for weak recovery.*

- ▶ **Our goal:** Retrieve and extend this result using unified **Bayesian MCMC coupling**.

02

The Bayesian Framework & Coupling






The General Bayesian Model

Definition

1. **Nodes, Edges, Communities:** Nodes $V_n = \{1, \dots, n\}$, edges E_n , communities $\mathcal{L}_K = \{0, \dots, K-1\}$, and set of configurations $\sigma \in \mathcal{C}_n = (\mathcal{L}_K)^{V_n}$.
2. **Features:** $X_n \in \mathcal{X}_n$ according to distribution ν_n .
3. **Prior on configurations:** $\mu_{0,n}(d\sigma \mid x)$ over $\mathcal{C}_n$.
4. **(Noisy) channel:** $Q_n(dw \mid x, \sigma) = \prod_{e \in E_n} Q_{e,n}(dw_e \mid x_i, x_j, \sigma_i, \sigma_j)$.

The joint law is $\mathbb{P}_n(dx, d\sigma, dw) = \nu_n(dx) \mu_{0,n}(d\sigma \mid x) Q_n(dw \mid x, \sigma)$.





Posterior Resampling (Nishimori Identity)

Replacing Truth by a Replica

- ▶ Comparing the algorithm's output with the hidden truth Σ_n is equivalent to comparing it with a posterior replica $\Sigma_n^{(1)}$.

Theorem (Posterior Resampling / Nishimori Identity [5])

Let τ_n be any community detection algorithm. Let $\Sigma_n^{(1)} \sim \mu_{n, X_n, W_n}$ be a replica drawn independently from the posterior. For any bounded test function Ψ_n :

$$\mathbb{E} [\Psi_n(X_n, W_n, \Sigma_n, \tau_n)] = \mathbb{E} [\Psi_n(X_n, W_n, \Sigma_n^{(1)}, \tau_n)] .$$

Corollary (Overlap Identity)

For any algorithm τ_n and threshold $s \in [0, 1]$:


$$\mathbb{P} [\text{ov}_n(\Sigma_n, \tau_n) \geq s] = \mathbb{P} [\text{ov}_n(\Sigma_n^{(1)}, \tau_n) \geq s] .$$



Invariant Coupling

Markov Chain Dynamics as a Source of Coupling

- ▶ We construct the posterior replica $\Sigma^{(1)}$ starting from the hidden truth Σ .
- ▶ Running **exactly one step** of a posterior-invariant MARKOV kernel M initialized at Σ yields a valid replica $\Sigma^{(1)} \sim M(\Sigma, \cdot)$.



03


MCMC Dynamics & Performance Bounds



Single-Spin Dynamics & Geometric Bottleneck


Single-Spin Update Limitations

- **Geometry makes updates slow:** Single-spin dynamics (GLAUBER / METROPOLIS-HASTINGS [6]) update one node at a time.
- **Contradictory Clusters:** Two internally valid (highly satisfied) but mutually contradictory clusters are locked.
- **High Energy Barrier:** Flipping a single boundary node violates internal edges, creating a local energy barrier.
- **Conclusion:** Need for geometry $\implies$ **Cluster Dynamics** (updating whole regions).





One Step of Signed Swendsen–Wang [7, 8]



- Edges in bold represent satisfied interactions that are randomly frozen.
- Clusters are then flipped independently.



Swendsen–Wang Signed Cluster Dynamics

Swendsen–Wang (SW) Signed Cluster Dynamics [7, 8]

- ▶ Update whole clusters of vertices simultaneously to speed up mixing.
- ▶ **Freezing Step:** For each edge $e = \{i, j\}$:
 - If e is satisfied ($W_e > 0$ and $\sigma_i = \sigma_j$, or $W_e < 0$ and $\sigma_i \neq \sigma_j$), freeze it with probability:

$$p_e^{\text{frozen}} = 1 - e^{-|W_e|}.$$

- If e is unsatisfied, do not freeze it ($p_e^{\text{frozen}} = 0$).
- ▶ **Recoloring Step:** Flip each frozen cluster independently with probability $1/2$.



High-Order Triangular Dynamics

The Frustration Obstruction

- ▶ **Frustration:** $\nexists \sigma$ satisfying all edges (conflicting constraints).
- ▶ **Excessive Freezing:** In frustrated configurations, standard SW satisfies and freezes too many edges independently, blocking/freezing the MCMC dynamics.
- ▶ **Higher-Order Correlation:** Group interactions into **triangles** to correlate freezing and **freeze** as little as possible.

Triangular Cluster Dynamics [9]

- ▶ **Energy Decomposition:** Distribute edge energy U over triangles: $U(\sigma) = \sum_{t \in \mathcal{T}} U_t(\sigma)$.
- ▶ **Freezing Mechanism:** Apply a joint freezing mechanism on the triangles (rather than independent edges) to keep cluster sizes active.



Triangular Dynamics Freezing Rules

$\triangle$ Attractive Triangle (Isotropic)					
Configuration	$\emptyset$	$\{ab\}$	$\{ac\}$	$\{bc\}$	$\{ab, ac, bc\}$
$(+, +, +)$	b_u	0	0	0	a_u
Other	1	0	0	0	0

$\triangleleft \triangle$ Repulsive Triangle (Isotropic)					
Configuration	$\emptyset$	$\{ab\}$	$\{ac\}$	$\{bc\}$	$\{ab, ac, bc\}$
$(+, +, +)$	1	0	0	0	0
$(+, +, -)$	b_u	0	$a_u/2$	$a_u/2$	0
$(+, -, +)$	b_u	$a_u/2$	0	$a_u/2$	0
$(+, -, -)$	b_u	$a_u/2$	$a_u/2$	0	0

Table: Freezing rules for triangular dynamics ($a_u = 1 - e^{-2u}$, $b_u = e^{-2u}$), where u is the energy transferred by each edge.






General Swendsen–Wang Cluster Dynamics

General Framework [10]

Suppose the Hamiltonian decomposes as: $U(\sigma) = U_0(\sigma) + \sum_{b \in \mathcal{B}} U_b(\sigma)$, where each $b \in \mathcal{B}$ is a **bond** (e.g. edge, triangle, clique).

1. **Freezing Step:** For each bond $b \in \mathcal{B}$, freeze a sub-bond $\omega_b \subseteq b$ with probability $P_b^\sigma(\omega_b)$.
2. **Recoloring Step:** Recolor each frozen cluster independently according to a distribution preserving the prior $\mu_0(\sigma) \propto e^{-U_0(\sigma)}$.





General Swendsen–Wang Cluster Dynamics

Definition (Local Reversibility)

For two configurations σ, σ' and a bond b , we define the compatibility probability:

$$P_b^{\sigma \rightarrow \sigma'}(\omega_b) = \begin{cases} P_b^\sigma(\omega_b), & \text{if } P_b^{\sigma'}(\omega_b) > 0, \\ 0, & \text{otherwise.} \end{cases}$$

The freezing rule is **locally reversible** if for all $b, \omega_b, \sigma, \sigma'$:

$$P_b^{\sigma \rightarrow \sigma'}(\omega_b) = e^{-(U_b(\sigma') - U_b(\sigma))} P_b^{\sigma' \rightarrow \sigma}(\omega_b)$$

Theorem (Gibbs Invariance [10])

If freezing choices are independent, locally reversible, and $P_b^\sigma(\emptyset) > 0$ for all b, σ :

1. The global cluster dynamics is reversible with respect to $\mu(\sigma) \propto e^{-U(\sigma)}$.
2. The condition $P_b^\sigma(\emptyset) > 0$ ensures the **irreducibility** of the MARKOV chain.

Percolation as a Necessary Condition for Recovery

Theorem (Percolation Bound on Overlap [10])

Under the uniform prior μ_0 , if frozen clusters are recolored independently and uniformly, then for the one-step posterior replica $\Sigma_n^{(1)}$, any algorithm τ_n , and any $\eta > 0$:

$$\lim_{n \rightarrow \infty} \mathbb{P} \left[\text{ov}_n(\Sigma_n^{(1)}, \tau_n) \geq \frac{1}{K} + \frac{K-1}{K}(\theta^{\max} + \eta) \right] = 0$$

where $\theta^{\max}$ is the asymptotic fraction of nodes in giant frozen components.

Corollary (Percolation is Necessary)

*If there is no percolation ($\theta^{\max} = 0$), the maximum overlap is bounded by $1/K$ (random guess): weak recovery is impossible. Hence, **percolation is a necessary condition for weak recovery**.*

Key Intuition

Small clusters are recolored independently, washing out their contribution. Only giant (macroscopic) components can carry global community information.



Coupling & the Sankararaman–Baccelli Bound

One Step of Classical Edge-by-Edge Swendsen–Wang

Let's run one step of standard SWENDSEN–WANG (edges) initialized at the hidden truth Σ . For any edge e , the probability that it is frozen is:

► **Case 1** ($\Sigma_i = \Sigma_j$):

$$\mathbb{P}(\text{frozen}) = f_{\text{in}}(r_e) \left(1 - e^{-|W_e|}\right) = f_{\text{in}}(r_e) \left(1 - \frac{f_{\text{out}}(r_e)}{f_{\text{in}}(r_e)}\right) = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

► **Case 2** ($\Sigma_i \neq \Sigma_j$):

$$\mathbb{P}(\text{frozen}) = (1 - f_{\text{out}}(r_e)) \left(1 - e^{-|W_e|}\right) = (1 - f_{\text{out}}(r_e)) \left(1 - \frac{1 - f_{\text{in}}(r_e)}{1 - f_{\text{out}}(r_e)}\right) = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

Corollary (Retrieving the S&B Bound [1])

*Under the classical edge SW dynamics, the frozen graph is **exactly** an independent bond percolation with edge probability:*

$$p_e = f_{\text{in}}(r_e) - f_{\text{out}}(r_e).$$



Strict Improvement of the S&B bound (on the triangular grid)

Homogeneous Model ($f_{\text{in}} = p$, $f_{\text{out}} = 1 - p$)

On the triangular grid, we select a family of independent white triangles (like a chessboard).

- **Edge Dynamics (Standard SW):** Frozen graph is independent edge percolation with parameter $p_{\text{edge}} = 2p - 1$. Critical threshold:

$$p_c^{\text{edge}} = \frac{1}{2} + \sin\left(\frac{\pi}{18}\right) \approx 0.673\dots$$

- **Triangular Dynamics:** Frozen graph yields correlated triangle percolation. Critical threshold:

$$p_c^{\triangle} = \frac{7}{4} - \frac{\sqrt{17}}{4} \approx 0.7192\dots$$

Corollary (Strictly Better Impossibility Bound [10])

Since $p_c^{\text{edge}} < p_c^{\triangle}$, for any $p \in [0.673\dots, 0.7192\dots)$, the standard S&B bound cannot conclude (edge dynamics percolate), whereas triangular dynamics do not, proving that weak recovery is impossible.



Empirical Results

04

Open Questions & Conclusion





Open Questions

Limits of the One-Replica Coupling

- ▶ Our Bayesian framework combined with MCMC cluster dynamics allows us to construct **one** replica $\Sigma_n^{(1)}$.
- ▶ This only yields a bound (a **necessary condition** for weak recovery).

Towards a Necessary and Sufficient Condition

- ▶ To achieve equality (a **necessary and sufficient condition**), a framework involving **two replicas** would be required.
- ▶ How can we define and analyze such a two-replica coupling under MCMC dynamics?

Towards a Practical Community Detection Algorithm

Limits of Direct Gibbs Sampling

- ▶ Empirically, recovering communities by sampling directly from the GIBBS distribution does not perform well.
- ▶ This is likely because the **mixing time** of the MARKOV chain is too large.

Estimating the Correlation Matrix

- ▶ Constructing an estimator of the correlation matrix:

$$\mathbb{P}[\mathbb{I}_{\{\sigma_i = \sigma_j\}}]$$

gives significantly better results.

- ▶ **Future work:** It would be highly relevant to theoretically study why this approach is much more robust.



Distance Adapted to Cluster Dynamics

Hamming Similarity on the Quotient Space

- ▶ The standard overlap is essentially a **HAMMING** similarity on the configuration space quotiented by label permutations:

$$\text{ov}_n(\sigma, \sigma') = \max_{\pi \in \mathfrak{S}_K} \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\sigma'_i = \pi(\sigma_i)\}}$$

Markov Transition-Adapted Distance

- ▶ We should consider a distance that better reflects the transitions of the **MARKOV** cluster dynamics.
- ▶ For classical **SWENDSEN-WANG** (edges), this corresponds to the **symmetric difference** (the sum of edge weights satisfied by exactly one of the configurations):
$$d_{\text{sym}}(\sigma, \sigma') = \sum_{e=\{i,j\}} |W_e| \left| \mathbb{I}_{\{\sigma \text{ satisfies } e\}} - \mathbb{I}_{\{\sigma' \text{ satisfies } e\}} \right|.$$
- ▶ This would also establish a link to the **spectral analysis of the graph** (e.g. via Laplacian eigenvalues or mixing rates).



Bibliography I

- [1] Abishek Sankararaman and François Baccelli. “Community Detection on Euclidean Random Graphs”. In: *Proceedings of the Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)* (2018), pp. 2181–2200. DOI: 10.1137/1.9781611975031.142.
- [2] François Baccelli and Bartłomiej Błaszczyszyn. *Stochastic Geometry and Wireless Networks, Volume I - Theory*. Vol. 1. Foundations and Trends in Networking. NoW Publishers, 2009, p. 150. DOI: 10.1561/13000000006.
- [3] Jörg Reichardt and Stefan Bornholdt. “Statistical mechanics of community detection”. In: *Phys. Rev. E* 74 (1 July 2006), p. 016110. DOI: 10.1103/PhysRevE.74.016110.
- [4] Mathew Penrose. *Random Geometric Graphs*. Vol. 5. Oxford Studies in Probability. Oxford University Press, 2003. DOI: 10.1093/acprof:oso/9780198506263.001.0001.
- [5] Hidetoshi Nishimori. “Exact results and critical properties of the Ising model with competing interactions”. In: *Journal of Physics C: Solid State Physics* 13.21 (1980), pp. 4071–4076. DOI: 10.1088/0022-3719/13/21/012.
- [6] Persi Diaconis. “The Markov chain Monte Carlo revolution”. In: *Bulletin of the American Mathematical Society* 46.2 (2009), pp. 179–205. DOI: 10.1090/S0273-0979-08-01238-X.

Bibliography II

- [7] Robert H. Swendsen and Jian-Sheng Wang. “Nonuniversal critical dynamics in Monte Carlo simulations”. In: *Phys. Rev. Lett.* 58 (2 Jan. 1987), pp. 86–88. DOI: [10.1103/PhysRevLett.58.86](https://doi.org/10.1103/PhysRevLett.58.86).
- [8] Robert G. Edwards and Alan D. Sokal. “Generalization of the Fortuin-Kasteleyn-Swendsen-Wang representation and Monte Carlo algorithm”. In: *Phys. Rev. D* 38 (6 Sept. 1988), pp. 2009–2012. DOI: [10.1103/PhysRevD.38.2009](https://doi.org/10.1103/PhysRevD.38.2009).
- [9] Louis Hauseux, Nahuel Soprano-Loto, and Konstantin Avrachenkov. “Higher-order Monte Carlo cluster dynamics for community detection in Euclidean graphs”. In: *Asilomar Conference on Signals, Systems, and Computers 2025*. Pacific Grove (CA), United States, Oct. 2025. URL: <https://inria.hal.science/hal-05267074>.
- [10] Louis Hauseux, Nahuel Soprano-Loto, and Konstantin Avrachenkov. *A Bayesian Framework for Community Detection on Signed Graphs. Tighter Bounds with Higher-Order Interaction Markov Chain Monte Carlo Cluster Dynamics*. Article en cours de rédaction. 2026+.

Merci.

