

Community Detection on Signed Graphs: A Bayesian MCMC Coupling Approach



Seminar of the MathNet team (Inria Paris)

Louis HAUSEUX (Inria Sophia Antipolis, superv. by K. AVRACHENKOV & J. ZERUBIA)

Joint work with Nahuel SOPRANO-LOTO (Inria Paris)
and Konstantin AVRACHENKOV (Inria Sophia Antipolis)

Inria Paris, Room Paul ERDÖS (A115) – 2 p.m. Monday 15 June 2026



Outline

1. Motivation & the Sankararaman–Baccelli Model
2. The Bayesian Framework & Coupling
3. MCMC Dynamics & Performance Bounds
4. Open Questions & Conclusion

01

Motivation & the Sankararaman–Bacelli Model



Haplotype Assembly

1. Biological Setup: Two Homologous Chromosomes (Haplotypes)

Humans are diploid: they have two copies of each chromosome (homologous chromosomes).

Haplotype 1
(Truth $\Sigma = +$)


Haplotype 2
(Truth $\Sigma = -$)



Haplotype Assembly

2. High-Throughput Sequencing: DNA Reads

Sequencing generates short, overlapping fragments called **reads** from both chromosomes.






Haplotype Assembly

3. Modeling as a Signed Graph

Nodes = reads Edges = overlaps

Goal: Reconstruct the true haplotype assignments Σ from this observed signed graph.





Motivation: Signed Community Detection

Community Detection on Signed Graphs ($K = 2$)

- ▶ We aim to partition nodes into $K = 2$ communities: $\Sigma_i \in \{+1, -1\}$.
- ▶ Observed links represent **attractions** (solid lines) and **repulsions** (dashed lines).
- ▶ Under the hood, this is a **Stochastic Block Model** (SBM) setup.

The Geometric Challenge

- ▶ Nodes are embedded in space (e.g., genomic positions on $\mathbb{Z}$ or $\mathbb{R}^d$).
- ▶ Edge probabilities depend directly on the distance $r = \|x_i - x_j\|$:

$$\mathbb{P}(\text{attractive edge} \mid \text{same community}) = f_{\text{in}}(r)$$

$$\mathbb{P}(\text{repulsive edge} \mid \text{different communities}) = f_{\text{out}}(r)$$

- ▶ **All the difficulty (and the interest!) comes from this geometric dependency.**



The Sankararaman–Bacelli Model (GSBM)

Model Definition [1]

- ▶ **Spatial Nodes** V_n : Hom. POISSON PP $\mathcal{X}$ in $\mathbb{R}^d$ [2] restrict. to a n -volume window.
- ▶ **Latent Communities**: Uniform assignments $\Sigma_i \in \{+1, -1\}$ ($K = 2$).
- ▶ **Geometric Overlaps**: Observed connection $e = \{i, j\} \in H$ with probability:

$$\mathbb{P}(e \in H \mid X, \Sigma) = \begin{cases} f_{\text{in}}(r_e) & \text{if } \Sigma_i = \Sigma_j \text{ (attraction),} \\ f_{\text{out}}(r_e) & \text{if } \Sigma_i \neq \Sigma_j \text{ (repulsion),} \end{cases}$$

with distance $r_e = \|X_i - X_j\|$ and connecting functions $f_{\text{in}}(r) \geq f_{\text{out}}(r)$.

Goal: Weak Recovery

- ▶ Find an algorithm τ (independent of Σ) reconstructing assignments:

$$\lim_{n \rightarrow \infty} \mathbb{E}[\text{ov}_n(\Sigma, \tau)] \geq \frac{1}{2} + \eta \quad \text{for some } \eta > 0,$$

using the **Community Overlap**: $\text{ov}_n(\Sigma, \tau) = \max_{\pi \in \mathfrak{S}_2} \frac{1}{n} \sum_{i=1}^n \mathbb{I}_{\{\tau_i = \pi(\Sigma_i)\}}.$



Gibbs Model on Signed Graphs

Attractive and Repulsive Edges

We partition the potential edges $E = F \dot{\cup} \tilde{F}$:

- ▶ F : **Attractive** (ferromagnetic) edges, with weight $W_e > 0$
- ▶ $\tilde{F}$: **Repulsive** (antiferromagnetic) edges, with weight $W_e < 0$

Signed Gibbs Measure [3]

The probability of a label configuration $\sigma \in \{+1, -1\}^{V_n}$ is $\mu(\sigma) = \frac{1}{Z} e^{-U(\sigma)}$, where the energy $U(\sigma)$ penalizes unsatisfied edges:

$$U(\sigma) = \sum_{e=\{i,j\} \in F} W_e \mathbb{I}_{\{\sigma_i \neq \sigma_j\}} + \sum_{e=\{i,j\} \in \tilde{F}} |W_e| \mathbb{I}_{\{\sigma_i = \sigma_j\}}$$

The magnitudes $|W_e|$ determine the strength of each constraint.



From GSBM to the Gibbs Posterior

Conditioning on the Point Cloud

- ▶ With fixed spatial positions X , all pairwise distances $r_e = \|X_i - X_j\|$ are known.
- ▶ Under a uniform prior on Σ , the **posterior** distribution given the observed graph H is:

$$\mathbb{P}(\Sigma = \sigma \mid H, X) \propto \mathbb{P}(H \mid X, \sigma)$$

Log-Likelihood as Signed Gibbs Energy

- ▶ Choosing weights $W_e(H)$ as below:

$$W_e(H) = \begin{cases} \log \left(\frac{f_{\text{in}}(r_e)}{f_{\text{out}}(r_e)} \right) > 0 & \text{if } e \in H \text{ (attractive } \in F), \\ -\log \left(\frac{1 - f_{\text{out}}(r_e)}{1 - f_{\text{in}}(r_e)} \right) < 0 & \text{if } e \notin H \text{ (repulsive } \in \tilde{F}). \end{cases}$$

- ▶ ... yields the exact correspondence:

$$\mathbb{P}(H \mid X, \sigma) \propto e^{-U(\sigma)}$$



From GSBM to the Gibbs Posterior

Conditioning on the Point Cloud

- ▶ With fixed spatial positions X , all pairwise distances $r_e = \|X_i - X_j\|$ are known.
- ▶ Under a uniform prior on Σ , the **posterior** distribution given the observed graph H is:

$$\mathbb{P}(\Sigma = \sigma \mid H, X) = \mu(\sigma) \propto \mathbb{P}(H \mid X, \sigma)$$

Log-Likelihood as Signed Gibbs Energy

- ▶ Choosing weights $W_e(H)$ as below:

$$W_e(H) = \begin{cases} \log \left(\frac{f_{\text{in}}(r_e)}{f_{\text{out}}(r_e)} \right) > 0 & \text{if } e \in H \text{ (attractive } \in F), \\ -\log \left(\frac{1 - f_{\text{out}}(r_e)}{1 - f_{\text{in}}(r_e)} \right) < 0 & \text{if } e \notin H \text{ (repulsive } \in \tilde{F}). \end{cases}$$

- ▶ ... yields the exact correspondence:

$$\mathbb{P}(H \mid X, \sigma) \propto e^{-U(\sigma)}$$



The Percolation Obstruction

Necessary Condition for Weak Recovery [1]

- ▶ Fragmented graphs (small, isolated components) prevent global reconstruction.
- ▶ Global information flow requires the graph to **percolate** [4] (possess a giant component).

Theorem (Sankararaman & Baccelli [1])

Independent bond percolation with edge probability

$$p_e = f_{\text{in}}(r_e) - f_{\text{out}}(r_e)$$

*yielding a giant component is a **necessary condition** for weak recovery.*

- ▶ **Our goal:** Retrieve and extend this result using unified **Bayesian MCMC coupling**.

02

The Bayesian Framework & Coupling






The General Bayesian Model

Definition

1. **Nodes, Edges, Communities:** Nodes $V_n = \{1, \dots, n\}$, edges E_n , communities $\mathcal{L}_K = \{0, \dots, K-1\}$, and set of configurations $\sigma \in \mathcal{C}_n = (\mathcal{L}_K)^{V_n}$.
2. **Features:** $X_n \in \mathcal{X}_n$ according to distribution ν_n .
3. **Prior on configurations:** $\mu_{0,n}(d\sigma \mid x)$ over $\mathcal{C}_n$.
4. **(Noisy) channel:** $Q_n(dw \mid x, \sigma) = \prod_{e \in E_n} Q_{e,n}(dw_e \mid x_i, x_j, \sigma_i, \sigma_j)$.

The joint law is $\mathbb{P}_n(dx, d\sigma, dw) = \nu_n(dx) \mu_{0,n}(d\sigma \mid x) Q_n(dw \mid x, \sigma)$.





Posterior Resampling (Nishimori Identity)

Replacing Truth by a Replica

- ▶ Comparing the algorithm's output with the hidden truth Σ_n is equivalent to comparing it with a posterior replica $\Sigma_n^{(1)}$.

Theorem (Posterior Resampling / Nishimori Identity [5])

Let τ_n be any community detection algorithm. Let $\Sigma_n^{(1)} \sim \mu_{n, X_n, W_n}$ be a replica drawn independently from the posterior. For any bounded test function Ψ_n :

$$\mathbb{E} [\Psi_n(X_n, W_n, \Sigma_n, \tau_n)] = \mathbb{E} [\Psi_n(X_n, W_n, \Sigma_n^{(1)}, \tau_n)] .$$

Corollary

Overlap Identity For any algorithm τ_n and threshold $s \in [0, 1]$:


$$\mathbb{P} [\text{ov}_n(\Sigma_n, \tau_n) \geq s] = \mathbb{P} [\text{ov}_n(\Sigma_n^{(1)}, \tau_n) \geq s] .$$



Invariant Coupling

Markov Chain Dynamics as a Source of Coupling

- ▶ We construct the posterior replica $\Sigma^{(1)}$ starting from the hidden truth Σ .
- ▶ Running **exactly one step** of a posterior-invariant MARKOV kernel M initialized at Σ yields a valid replica $\Sigma^{(1)} \sim M(\Sigma, \cdot)$.



03


MCMC Dynamics & Performance Bounds



Single-Spin Dynamics & Geometric Bottleneck


Single-Spin Update Limitations

- **Geometry makes updates slow:** Single-spin dynamics (GLAUBER / METROPOLIS-HASTINGS [6]) update one node at a time.
- **Contradictory Clusters:** Two internally valid (highly satisfied) but mutually contradictory clusters are locked.
- **High Energy Barrier:** Flipping a single boundary node violates internal edges, creating a local energy barrier.
- **Conclusion:** Need for geometry $\implies$ **Cluster Dynamics** (updating whole regions).





One Step of Signed Swendsen–Wang [7, 8]



- Edges in bold represent satisfied interactions that are randomly frozen.
- Clusters are then flipped independently.



Swendsen–Wang Signed Cluster Dynamics

Swendsen–Wang (SW) Signed Cluster Dynamics [7, 8]

- ▶ Update whole clusters of vertices simultaneously to speed up mixing.
- ▶ **Freezing Step:** For each edge $e = \{i, j\}$:
 - If e is satisfied ($W_e > 0$ and $\sigma_i = \sigma_j$, or $W_e < 0$ and $\sigma_i \neq \sigma_j$), freeze it with probability:

$$p_e^{\text{gel}} = 1 - e^{-|W_e|}.$$

- If e is unsatisfied, do not freeze it ($p_e^{\text{gel}} = 0$).
- ▶ **Recoloring Step:** Flip each frozen cluster independently with probability $1/2$.



High-Order Triangular Dynamics

The Frustration Obstruction

- ▶ **Frustration:** Cycles with an odd number of negative interactions (conflicting constraints).
- ▶ **Excessive Freezing:** In frustrated graphs, standard SW satisfies and freezes too many edges independently, blocking/freezing the MCMC dynamics.
- ▶ **Higher-Order Idea:** Group interactions into **triangles** to correlate the freezing mechanism and **freeze as little as possible**.

Triangular Cluster Dynamics

- ▶ **Attractive Triangle** (+ + + or - - -): Freeze all 3 edges (prob $a_u = 1 - e^{-2u}$) or none (prob $b_u = e^{-2u}$).
- ▶ **Repulsive Triangle** (2 satisfied, 1 unsatisfied): Freeze **exactly one** of the two satisfied edges (prob $a_u/2$). Minimal freezing keeps the dynamics active.



Triangular Dynamics Freezing Rules

$\triangle$ Attractive Triangle (Isotropic)					
Configuration	$\emptyset$	$\{ab\}$	$\{ac\}$	$\{bc\}$	$\{ab, ac, bc\}$
$(+, +, +)$	b_u	0	0	0	a_u
Other	1	0	0	0	0

$\triangleleft \triangle$ Repulsive Triangle (Isotropic)					
Configuration	$\emptyset$	$\{ab\}$	$\{ac\}$	$\{bc\}$	$\{ab, ac, bc\}$
$(+, +, +)$	1	0	0	0	0
$(+, +, -)$	b_u	0	$a_u/2$	$a_u/2$	0

Table: Freezing rules for triangular dynamics ($a_u = 1 - e^{-2u}$, $b_u = e^{-2u}$). By construction, these rules satisfy local reversibility and leave the posterior distribution invariant.

Percolation as a Necessary Condition for Recovery

Theorem (Percolation Bound on Overlap [9])

Under the uniform community prior μ_0 (baseline overlap $1/K$), if the frozen clusters are recolored independently and uniformly, then for any algorithm τ_n and any $\eta > 0$:

$$\lim_{n \rightarrow \infty} \mathbb{P} \left[\text{ov}_n(\Sigma_n, \tau_n) \geq \frac{1}{K} + \frac{K-1}{K}(\theta^{\max} + \eta) \right] = 0$$

where $\theta^{\max}$ is the asymptotic fraction of nodes in giant frozen components.

Corollary

Percolation is Necessary In the absence of percolation ($\theta^{\max} = 0$), the maximum possible overlap is bounded by $1/K$ (random guess):

$$\theta^{\max} = 0 \implies \text{Weak (partial) recovery is impossible.}$$

Hence, percolation of the cluster dynamics is a necessary condition for weak recovery.

Key Intuition

Small (finite) clusters are recolored independently, washing out their contribution to the global alignment. Only giant (macroscopic) components can carry global community information.



Coupling & the Sankararaman–Bacelli Bound

Gel Probability under the Generative Model

Let's run one step of SW starting from the hidden truth Σ . What is the probability that an edge e is frozen?

- **Case 1:** $\Sigma_i = \Sigma_j$. Edge is present with probability $f_{\text{in}}(r_e)$, and has weight $W_e = \log(f_{\text{in}}/f_{\text{out}}) > 0$.

$$\mathbb{P}(\text{frozen}) = f_{\text{in}}(r_e) \left(1 - e^{-|W_e|}\right) = f_{\text{in}}(r_e) \left(1 - \frac{f_{\text{out}}(r_e)}{f_{\text{in}}(r_e)}\right) = f_{\text{in}}(r_e) - f_{\text{out}}(r_e).$$

- **Case 2:** $\Sigma_i \neq \Sigma_j$. Edge is absent with probability $1 - f_{\text{out}}(r_e)$, and has weight $W_e = -\log((1 - f_{\text{out}})/(1 - f_{\text{in}})) < 0$.

$$\mathbb{P}(\text{frozen}) = (1 - f_{\text{out}}(r_e)) \left(1 - \frac{1 - f_{\text{in}}(r_e)}{1 - f_{\text{out}}(r_e)}\right) = f_{\text{in}}(r_e) - f_{\text{out}}(r_e).$$

Corollary

*Independent Bond Percolation Under the true generative model, the graph of frozen edges is **exactly** an independent bond percolation with edge probability:*

$$p_e = f_{\text{in}}(r_e) - f_{\text{out}}(r_e).$$

This coupling directly retrieves the necessary percolation condition of SODA 2018!



Strict Improvement on the Triangular Grid

Homogeneous Model ($f_{\text{in}} = p$, $f_{\text{out}} = 1 - p$)

On the triangular grid, we select a family of independent white triangles (like a chessboard).

- ▶ **Edge Dynamics (Standard SW):** Gelled graph is independent edge percolation with parameter $p_{\text{edge}} = 2p - 1$. Critical threshold: $p_c^{\text{edge}} = \frac{1}{2} + \sin(\frac{\pi}{18}) \approx 0.673$.
- ▶ **Triangular Dynamics:** Gelled graph yields correlated triangle percolation. Critical threshold: $p_c^{\Delta} = \frac{7}{4} - \frac{\sqrt{17}}{4} \approx 0.7192$.

Corollary

Strictly Better Impossibility Bound [10] Since $p_c^{\text{edge}} < p_c^{\Delta}$, for any


$$p \in [0.673; 0.7192[,$$

standard edge dynamics percole (no conclusion), whereas triangular dynamics do not. Hence, we prove that weak recovery is impossible in this wider range of parameters!




Comparison of Percolation Thresholds

Triangular Grid



White triangles used by
triangular dynamics

Weak Recovery Impossibility Bounds



04

Open Questions & Conclusion





Open Questions

1. Extension to Other Lattices and Higher Dimensions

- ▶ What are the thresholds on the KAGOME, FCC, HCP, or Hyper-cubic lattices?
- ▶ Can we define tetrahedral or higher-order simplex dynamics to obtain even tighter bounds?

2. Multi-Community case ($K > 2$)

- ▶ How does the cluster dynamics extend when the spins take K values?
- ▶ What is the role of Potts-like percolation in the coupling?

3. Sufficiency of the Percolation Condition

- ▶ The percolation threshold provides a necessary condition (impossibility).
- ▶ Is percolation also sufficient for the existence of a reconstruction algorithm?



Bibliography I

- [1] Abishek Sankararaman and François Baccelli. “Community Detection on Euclidean Random Graphs”. In: *Proceedings of the Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)* (2018), pp. 2181–2200. DOI: 10.1137/1.9781611975031.142.
- [2] François Baccelli and Bartłomiej Błaszczyszyn. *Stochastic Geometry and Wireless Networks, Volume I - Theory*. Vol. 1. Foundations and Trends in Networking. NoW Publishers, 2009, p. 150. DOI: 10.1561/13000000006.
- [3] Jörg Reichardt and Stefan Bornholdt. “Statistical mechanics of community detection”. In: *Phys. Rev. E* 74 (1 July 2006), p. 016110. DOI: 10.1103/PhysRevE.74.016110.
- [4] Mathew Penrose. *Random Geometric Graphs*. Vol. 5. Oxford Studies in Probability. Oxford University Press, 2003. DOI: 10.1093/acprof:oso/9780198506263.001.0001.
- [5] Hidetoshi Nishimori. “Exact results and critical properties of the Ising model with competing interactions”. In: *Journal of Physics C: Solid State Physics* 13.21 (1980), pp. 4071–4076. DOI: 10.1088/0022-3719/13/21/012.
- [6] Persi Diaconis. “The Markov chain Monte Carlo revolution”. In: *Bulletin of the American Mathematical Society* 46.2 (2009), pp. 179–205. DOI: 10.1090/S0273-0979-08-01238-X.



Bibliography II

- [7] Robert H. Swendsen and Jian-Sheng Wang. “Nonuniversal critical dynamics in Monte Carlo simulations”. In: *Phys. Rev. Lett.* 58 (2 Jan. 1987), pp. 86–88. DOI: [10.1103/PhysRevLett.58.86](https://doi.org/10.1103/PhysRevLett.58.86).
- [8] Robert G. Edwards and Alan D. Sokal. “Generalization of the Fortuin-Kasteleyn-Swendsen-Wang representation and Monte Carlo algorithm”. In: *Phys. Rev. D* 38 (6 Sept. 1988), pp. 2009–2012. DOI: [10.1103/PhysRevD.38.2009](https://doi.org/10.1103/PhysRevD.38.2009).
- [9] Louis Hauseux, Nahuel Soprano-Loto, and Konstantin Avrachenkov. *A Bayesian Framework for Community Detection on Signed Graphs. Tighter Bounds with Higher-Order Interaction Markov Chain Monte Carlo Cluster Dynamics*. Article en cours de rédaction. 2026+.
- [10] Louis Hauseux, Nahuel Soprano-Loto, and Konstantin Avrachenkov. “Higher-order Monte Carlo cluster dynamics for community detection in Euclidean graphs”. In: *Asilomar Conference on Signals, Systems, and Computers 2025*. Pacific Grove (CA), United States, Oct. 2025. URL: <https://inria.hal.science/hal-05267074>.

Merci.

